



Conditional token readout: all 192 inputs, 33 layers

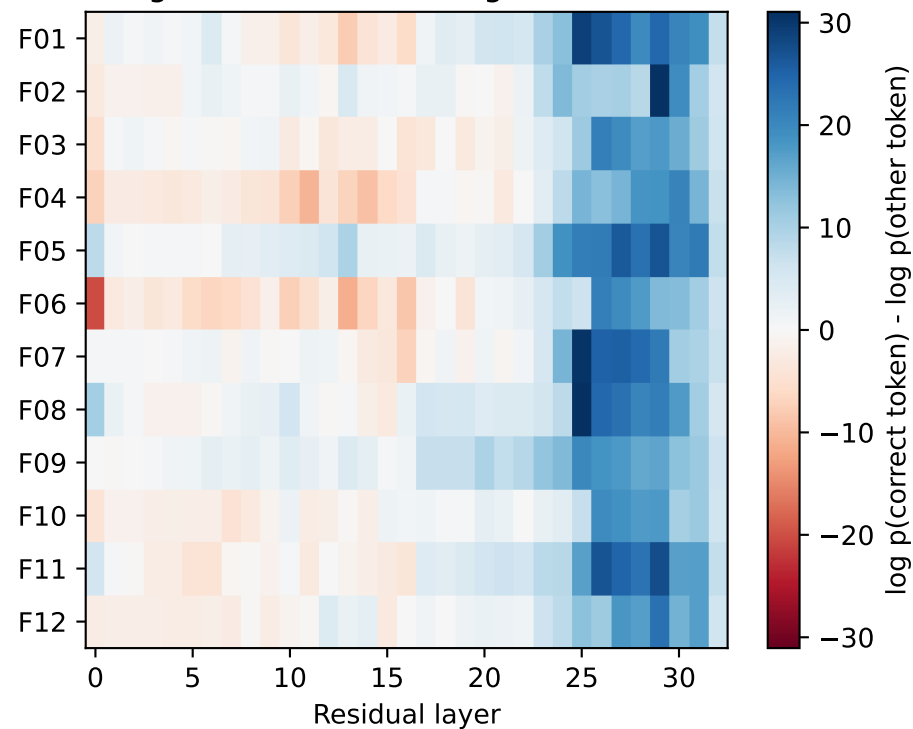
Target vs reversed



Target vs background/null



Target vs reversed background/both A



Each comparison uses its own shared completion prefix. These are not one normalized four-answer distribution.